




First token · GPU expert systems



First token · llama.cpp (separate scale)



Decode latency · lower is better



Entire batch throughput · higher is better

